

CODING TIME!

Imports

```
1 import os
2 import torch
3 import torch.distributed as dist
4 import torch.nn as nn
5 import torch.optim as optim
6 from torch.nn.parallel import DistributedData...
7 from torchvision import datasets, transforms,...
8 from torch.utils.data import DataLoader,...
```

These are standard PyTorch and torchvision libraries for:

- Building and training models
- Distributed training
- Using datasets and transformations

Distributed Setup

```
1 def setup():
2     dist.init_process_group(
3         backend='gloo',
4         init_method='env://',
5         world_size=int(os.environ['WORLD_SIZE']),
6         rank=int(os.environ['RANK'].split('-')[-1])
7     )
```

Explanation:

- Initializes the DDP process group.
- backend='gloo' for CPU or simple GPU communication.
- init_method='env://' uses environment variables.
- Extracts rank from strings like "ddp-trainer-0".

Distributed Setup

```
1 def setup():
2     dist.init_process_group(
3         backend='gloo',
4         init_method='env://',
5         world_size=int(os.environ['WORLD_SIZE']),
6         rank=int(os.environ['RANK'].split('-')[-1])
7     )
```

Explanation:

- Initializes the DDP process group.
- backend='gloo' for CPU or simple GPU communication.
- init_method='env://' uses environment variables.
- Extracts rank from strings like "ddp-trainer-0".

Cleanup

```
1 def cleanup():  
2     dist.destroy_process_group()
```

Releases resources and cleanly exits the process group after training.

Main Function (Start)

```
1 def main():  
2     setup()
```

Begins by initializing the distributed training environment.

Dataset and Transform

```
1 transform = transforms.Compose([  
2     transforms.ToTensor(),  
3     transforms.Normalize((0.5,), (0.5,))  
4 ])   
5   
6 dataset = datasets.CIFAR10(root='/srv/nfs/cifar10'...
```

Explanation:

- Converts images to tensors and normalizes them.
- CIFAR-10 dataset is downloaded to a shared NFS directory.

Distributed Sampler and DataLoader

```
1 sampler = DistributedSampler(dataset)
2 dataloader = DataLoader(dataset, sampler=sampler...
```

Explanation:

- DistributedSampler splits the dataset across processes.
- DataLoader uses the sampler to prevent data duplication.

Model and DDP Wrapping

```
1 model = models.resnet18(num_classes=10)
2 model = DDP(model)
```

Explanation:

- Loads ResNet18 with 10 output classes.
- Wraps the model with `DistributedDataParallel` for synced training.

Loss Function and Optimizer

```
1 loss_fn = nn.CrossEntropyLoss()  
2 optimizer = optim.Adam(model.parameters(), lr=0.001)
```

Explanation:

- Uses cross-entropy for classification.
- Optimizes using Adam with learning rate 0.001.

Training Loop

```
1 for epoch in range(3):
2     print("epoch", epoch)
3     ctr = 0
4     for images, labels in dataloader:
5         ctr += 1
6         print("hello123", ctr)
7         outputs = model(images)
8         loss = loss_fn(outputs, labels)
9         optimizer.zero_grad()
10        loss.backward()
11        optimizer.step()
12        if ctr % 10 == 0:
13            print(f"[RANK_{os.environ['RANK']}] ...")
14        print(f"[RANK {os.environ['RANK']}] ...")
```

Explanation:

- Trains for 3 epochs, printing progress every 10 steps.

Final Cleanup and Entry Point

```
1 cleanup()  
2  
3 if __name__ == "__main__":  
4     main()
```

Explanation:

- Cleans up the distributed environment at the end.
- Ensures `main()` only runs when script is executed directly.